

# Defeating Samsung KNOX with zero privilege

Di Shen a.k.a. Retme (@returnsme)

Keen Lab of Tencent



# whoami

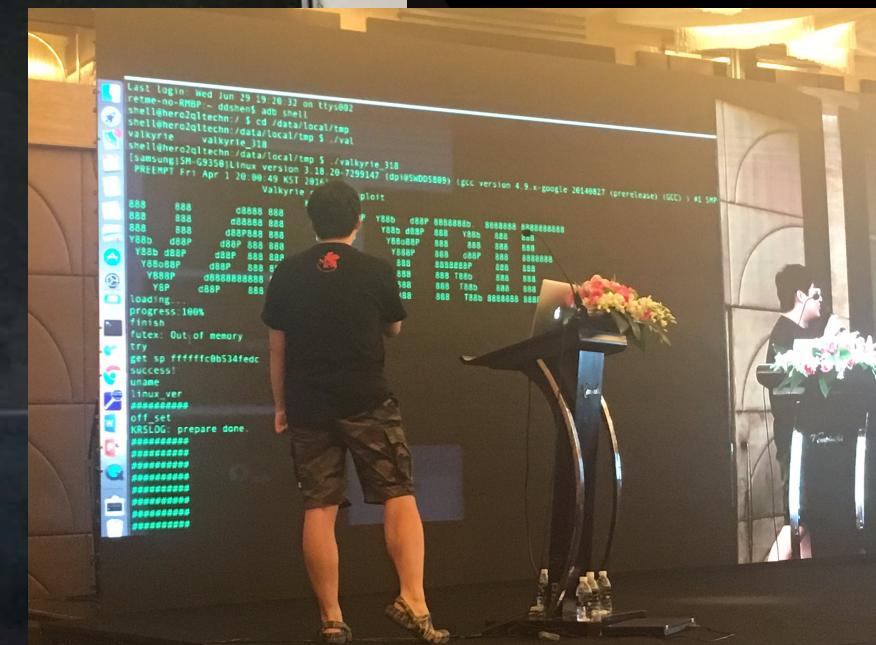
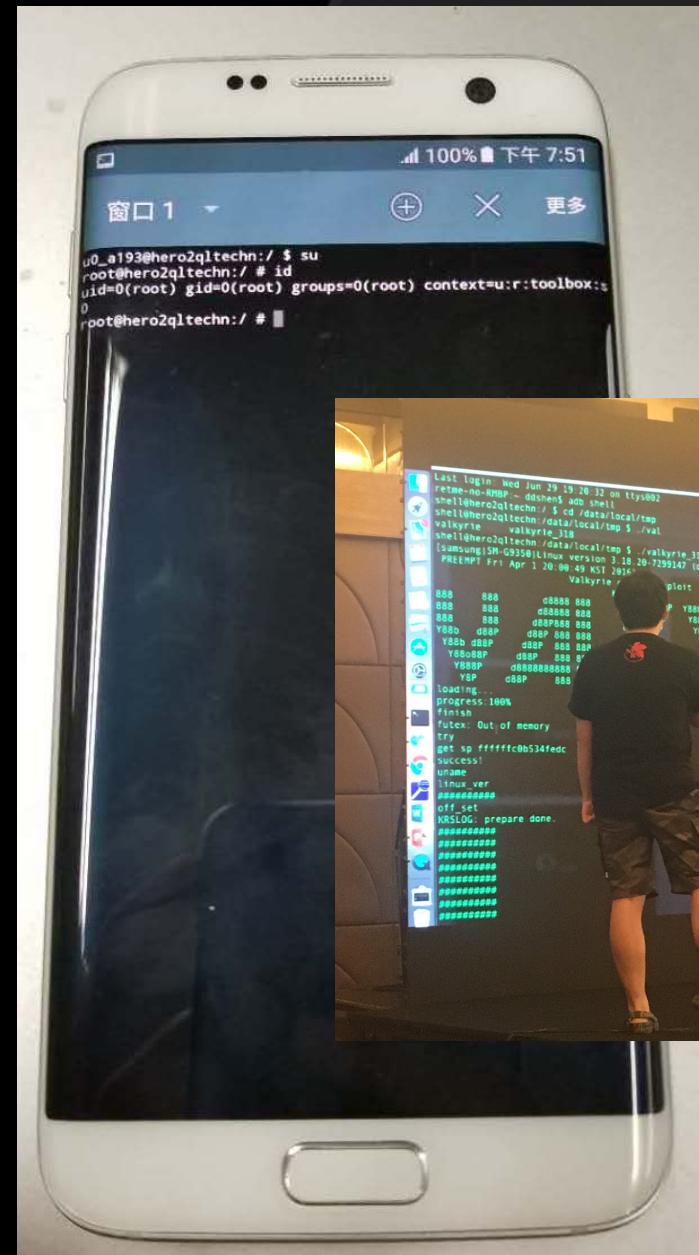
- Di Shen a.k.a. Retme (@returnsme)
- Member of Keen Lab
- Android Kernel vulnerability hunting and exploitation since 2014
- Aim: to make out universal rooting exploit for Android
- Trophy:
  - CVE-2016-6787 & CVE-2017-0403 (kernel/events/core.c)
    - Credit to discoveries and exploits
  - CVE-2015-1805 (fs/pipe.c)
    - First working exploit
  - CVE-2015-4421,4422
    - Kernel LPE and TrustZone code execution for Huawei Mate 7
  - Exploiting Wireless Extension for all common Wi-Fi chipsets (BHEU 16')
  - And more To Be Announced in the future

# Agenda

- Overview of KNOX 2.6
  - KASLR (Samsung's implementation)
  - Real-time kernel protection (RKP)
  - Data Flow Integrity (DFI)
  - SELinux enhancement
- Bypassing techniques
  - KASLR bypassing
  - DFI bypassing
  - SELinux bypassing
  - Gain root

# Target device

- Samsung Galaxy S7 edge
    - SM-G9350 (Hong Kong ver.)
    - Qualcomm-based
    - KNOX 2.6
  - Exploit chain was finished in June 2016
  - Demonstrated in July 1<sup>st</sup> 2016 at Shanghai



# Common LPE flow on Android

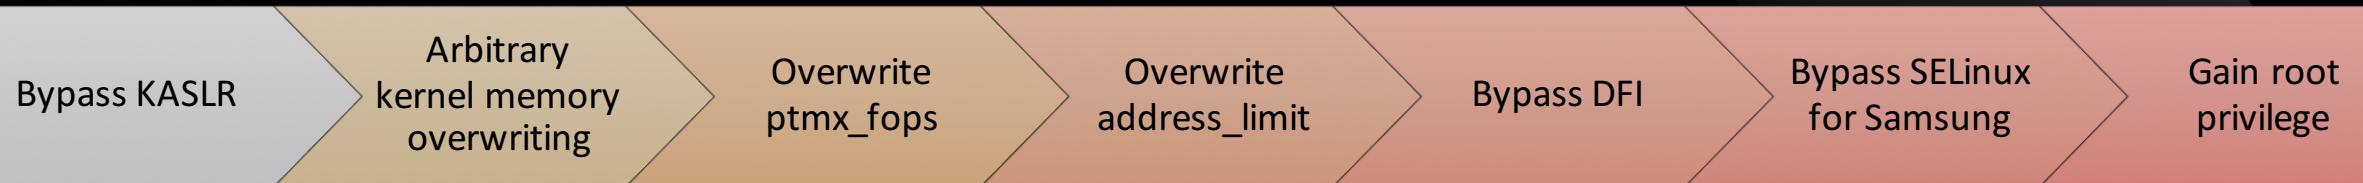
Arbitrary kernel memory  
overwriting

Overwrite  
ptmx\_fops

Overwrite  
address\_limit

Overwrite uid,  
security id, and  
selinux\_enforcing

# LPE flow on Galaxy S7 edge



# KASLR for Linux 3.18 - Initialization

- CONFIG\_RELOCATABLE\_KERNEL by Samsung
- The Random size is passed to kernel by loader
- X1,X2 are set upon kernel start up
  - X1: phycal offset X2: vitual text offset
  - Store to \_\_boot\_kernel\_offset
  - \_\_boot\_kernel\_offset[0] : physical address of kernel
  - \_\_boot\_kernel\_offset[1] : the actual load address
  - \_\_boot\_kernel\_offset[2] : TEXT\_OFFSET 0x8000

```
ENTRY(stext)
#endif CONFIG_RKP_CFP_ROPP
    /* Must intialize RRK to zero before any RET/BL */
    mov RRK, #0
#endif
#ifndef CONFIG_RKP_CFP_JOPP
    /* We need RRS to be loaded before we take our */
    /* load_function_entry_magic_number_before_reloc here */
#endif
#ifndef CONFIG_RELOCATABLE_KERNEL
    mov x22, x1                // x1=PHYS_OFFSET
    mov x19, x2                // x2=real TEXT_OFFSET
    adr x21, __boot_kernel_offset
    stp x1, x2, [x21]
#endif
```

# KASLR for Linux 3.18 - relocating

- `__relocate_kernel()` handles kernel relocating
  - Similar to a aarch64 linker in user space

```
#ifdef CONFIG_RELOCATABLE_KERNEL

#define R_AARCH64_RELATIVE 0x403
#define R_AARCH64_ABS64    0x101

__relocate_kernel:
    sub x23, [x19], #TEXT_OFFSET
    adrp    x8, __dynsym_start
    add x8, x8, :lo12:_dynsym_start //x8: start of symbol table
    adrp    x9, __reloc_start
    add x9, x9, :lo12:_reloc_start //x9: start of relocation table
    adrp    x10, __reloc_end
    add x10, x10, :lo12:_reloc_end //x10: end of relocation table
```

# KASLR for Linux 3.18 - .rela section

- `__relocate_kernel` handles kernel relocating
  - Similar to a aarch64 linker in user space
  - Relocation section ‘.rela’ at offset 0x1446600 contains 233903 entries:

Offset	Info	Type	Sym. Value	Sym. Name + Addend
fffffc000081698	000000000403	R_AARCH64_RELATIV		-3ffff7e968
fffffc0000816a0	000000000403	R_AARCH64_RELATIV		-3ffe5d1e20
fffffc000081798	000000000403	R_AARCH64_RELATIV		-3ffff7e868
...				
fffffc00008e000	000000000403	R_AARCH64_RELATIV		-3ffff546b8 # Begin of sys_call_table
...				
fffffc000f6f800	000600000101	R_AARCH64_ABS64	fffffc00080000 _text + 0	# Begin of kallsyms_addresses
...				
fffffc0013b1468	000000000403	R_AARCH64_RELATIV		-3ffec1afdd

# Bypassing KASLR

- Readable TIMA logs

```
shell@hero2qltechn:/proc $ ls -l | grep tz
.|shell@hero2qltechn:/proc $ ls -l | grep tima
·rw-r--r-- root      root          0 2016-05-14 16:52 tima_debug_log
·rw-r--r-- root      root          0 2016-05-14 16:52 tima_debug_rkp_log
·rw-r--r-- root      root          0 2016-05-12 19:44 tima_secure_rkp_log
shell@hero2qltechn:/proc $ █
```

# Kernel information leaking

Tencent

- Kernel pointer leaked in /proc/tima\_secure\_rkp\_log
- At 0x13B80 -> *init\_user\_ns*
- Real:0xFFFFFC001B0EFB8 Static:0xFFFFFC01A3AFB8
- KASLR offset = 0xD4000

		ROM:FFFFFC001A3AFB8	init_user_ns	DCB	1
1:3B00h:	0000000000000001	00 ROM:FFFFFC001A3AFBA		DCB	0
1:3B10h:	0000000000000000	00 ROM:FFFFFC001A3AFBB		DCB	0
1:3B20h:	0000000000000000	00 ROM:FFFFFC001A3AFBC		DCB	0
1:3B30h:	0000003FFFFFFFFF	00 ROM:FFFFFC001A3AFBD		DCB	0
1:3B40h:	0000003FFFFFFFFF	00 ROM:FFFFFC001A3AFBE		DCB	0
1:3B50h:	0000000000000000	00 ROM:FFFFFC001A3AFBF		DCB	0
1:3B60h:	0000000000000000	00 ROM:FFFFFC001A3AFCO		DCB	0
1:3B70h:	FFFFFC01C614400	00 ROM:FFFFFC001A3AFC1		DCB	0
1:3B80h:	FFFFFC001B0EFB8	00 ROM:FFFFFC001A3AFC2		DCB	0
1:3B90h:	0000000000000000	0000000000000000	.....	DCB	0
1:3BA0h:	FFFFFC0E8AE0100	FFFC0E783A400	ÿÿÿÀ.aD.ÿÿÿÀ.ºí`		
root@hero2qltechn:/ # cat /proc/kallsyms					
	fffffc00148cab8	R __ksymtab_init_user_ns	00000000		ÿÿÿÀ.Q ..
	fffffc001496883	r __kstrtab_init_user_ns	FFFFFFFFFF		ÿÿÿÀ.a<.ÿÿÿÿÿÿÿ
	fffffc001b0efb8	D init_user_ns	FFFFFFFFFF		ÿÿÿÿÿÿÿÿÿÿÿÿÿÿÿ

# Achieve arbitrary kernel mem overwriting

- By exploiting CVE-2016-6787
- Use-after-free due to race condition in perf subsystem
  - Moving group in sys\_perf\_event\_open() is not locked by mutex correctly
- Spray struct perf\_event\_context{}
  - Control code flow by refill ctx->pmu->pmu\_disable(X0)
- Another long story ☺

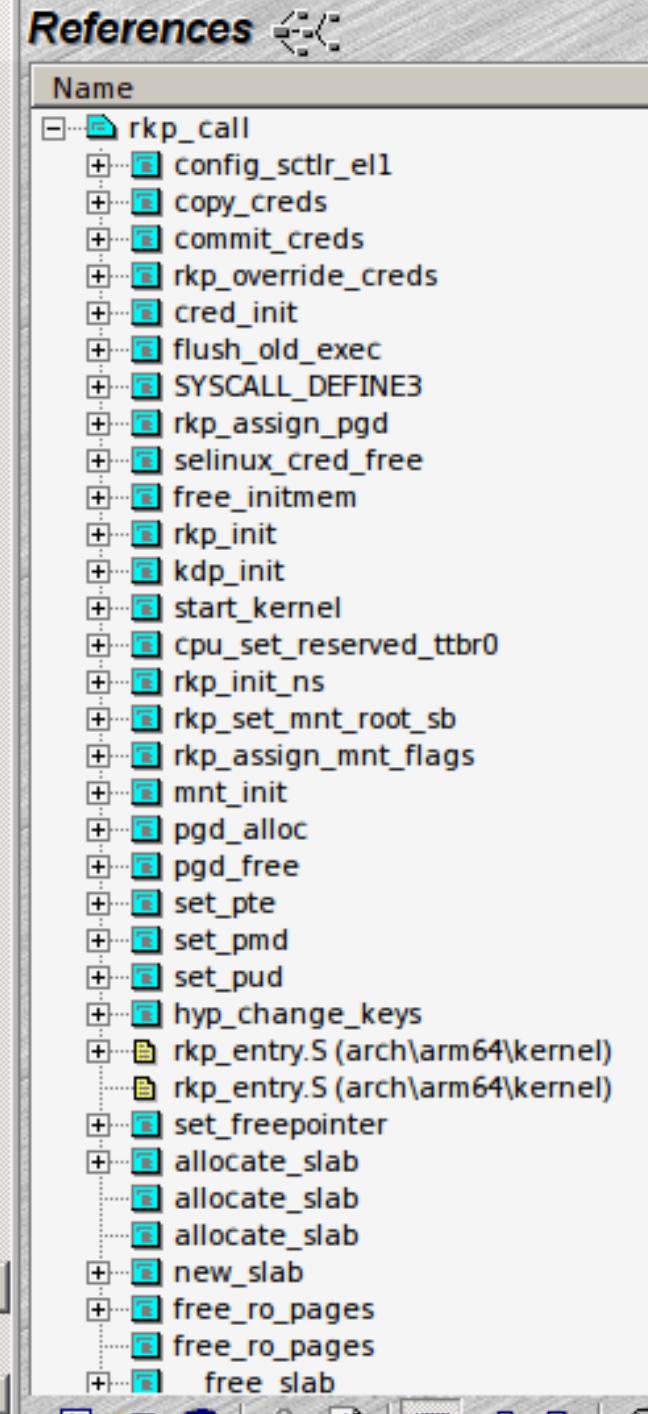
```
+     /*
+      * Take the group_leader's group_leader_mutex before observing
+      * anything in the group leader that leads to changes in ctx,
+      * many of which may be changing on another thread.
+      * In particular, we want to take this lock before deciding
+      * whether we need to move_group.
+      */
+     if (group_leader)
+         mutex_lock(&group_leader->group_leader_mutex);
+
+     if (pid != -1 && !(flags & PERF_FLAG_PID_CGROUP)) {
+         task = find_lively_task_by_vpid(pid);
+         if (IS_ERR(task)) {
@@ -7686,6 +7697,8 @@
             put_ctx(gctx);
         }
         mutex_unlock(&ctx->mutex);
+         if (group_leader)
+             mutex_unlock(&group_leader->group_leader_mutex);
```

# Real-time Kernel Protection

- Implemented in TrustZone or hypervisor
  - Depends on device model, for S7 edge (SM-G9350), it's TrustZone
- CONFIG\_TIMA\_RKP , CONFIG\_RKP\_KDP
- Targeted features via [samsungknox.com](http://samsungknox.com):
  - "completely prevents running unauthorized privileged code"
  - "prevents kernel data from being directly accessed by user processes"
  - "monitors some critical kernel data structures to verify that they are not exploited by attacks"

# rkp\_call()

- RKP call entry
  - Called by many critical kernel functions
    - SLAB allocation and de-allocation
    - Page Table operations
    - Copy/Override/Commit creds



# Kernel code protection

- Not exclusive features for KNOX 2.6
  - “config KERNEL\_TEXT\_RDONLY”
  - Data section not executable
  - Privileged eXecute Never (Pxn)
    - Kill ret2user and other ancient tricks
- New in KNOX 2.8?
  - Control flow protection

## PxnTable, bit[59]

Pxn limit for subsequent levels of lookup, see [Hierarchical control of instruction fetching](#) page D5-1786.

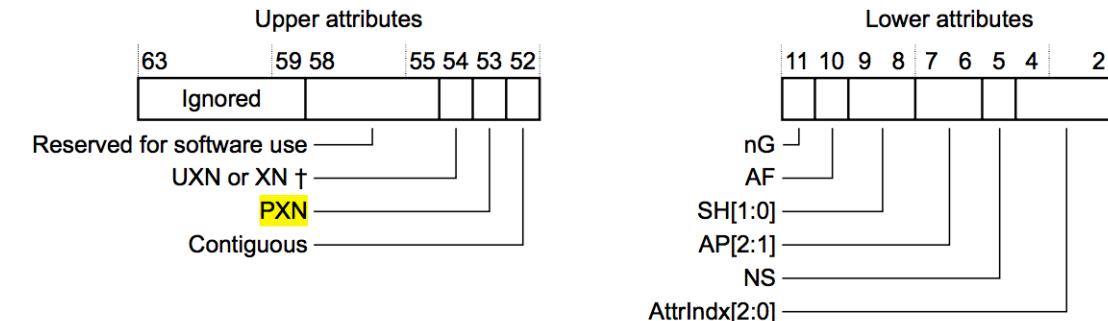
This bit is reserved, SBZ:

- In the EL2 translation regime.
- In the EL3 translation regime.

## Attribute fields in stage 1 VMSAv8-64 Block and Page descriptors

In Block and Page descriptors, the memory attributes are split into an upper block and a lower block, as shown in a stage 1 translation:

Attribute fields for VMSAv8-64 stage 1 Block and Page descriptors



# Kernel page and page table protection

- rkp\_call() handles :
  - allocations, de-allocations of page table
  - manipulations of page entries
- Neither kernel or user space can change attributes of protected pages unauthorizedly
- Related functions
  - pdg\_alloc/free()
  - set\_pte/pmd/pud()

```
static inline void set_pte(pte_t *ptep, pte_t pte)
{
#ifdef CONFIG_TIMA_RKP
    if (pte && rkp_is_pg_dbl_mapped((u64)(pte))) {
        panic("TIMA RKP : Double mapping Detected pte = %lx\n", pte);
        return;
    }
    if (rkp_is_pte_protected((u64)ptep)) {
        rkp_flush_cache((u64)ptep);
        rkp_call(RKP_PTE_SET, (unsigned long)ptep, pte);
        rkp_flush_cache((u64)ptep);
    } else {
        asm volatile(
            "mov x1, %0\n"
            "mov x2, %1\n"
            "str x2, [x1]\n"
            :
            : "r" (ptep), "r" (pte)
            : "x1", "x2", "memory" );
    }
#endif
    *ptep = pte;
#endif /* CONFIG_TIMA_RKP */
```

# Kernel data protection

- Based on read only pages
- Read-only global variables
  - RO after initialization
- RKP\_RO\_AREA located in page \_\_rkp\_ro\_start[]
  - struct cred init\_cred
  - struct task\_secrity\_struct init\_sec
  - struct security\_oprations security\_ops

```
#define RKP_RO_AREA __attribute__ ((section (".rkp.prot.page")))
extern int rkp_cred_enable;
extern char __rkp_ro_start[], __rkp_ro_end[];

extern struct cred init_cred;
```

# Kernel object protection

- Allocated in Read-only pages
  - Writable for hypervisor or TrustZone
- Protected Object type (name of its kmem\_cache):
  - cred\_jar\_ro : credential of processes
  - tsec\_jar: security context
  - vfsmnt\_cache: struct vfsmount
- Allocation, deallocation and overwriting routines will
  - call rpk\_call() to operate read-only objects
- Prevent kernel/user mode manipulating credentials, security context and mount namespace

# History: bypassing trick on S6

- Kernel Object protection had been applied on S6
- Could be bypassed by calling `rkp_override_creds()`
  - able to override current process's credentials via `rkp_call()` in secure world
- Not working on S7
  - S7 add more checking in secure world

- On S6, attacker can call this function to bypass previous kernel object protection

# Case study: rkp\_override\_creds()

```
#ifdef CONFIG_RKP_KDP
const struct cred *rkp_override_creds(struct cred **cnew)
#else
const struct cred *override_creds(const struct cred *new)
#endif /* CONFIG_RKP_KDP */
{
    const struct cred *old = current->cred;
#ifdef CONFIG_RKP_KDP
    struct cred *new = *cnew;
    struct cred *new_ro;
    volatile unsigned int rkp_use_count = rkp_get_usecount(new);
    void *use_cnt_ptr = NULL;
    void *tsec = NULL;
#endif /* CONFIG_RKP_KDP */

    kdebug("override_creds(%p[%d,%d])", new,
           atomic_read(&new->usage),
           read_cred_subscribers(new));

    validate_creds(old);
    validate_creds(new);
#ifdef CONFIG_RKP_KDP
    if(rkp_cred_enable) {
        cred_param_t cred_param;
        new_ro = kmem_cache_alloc(cred_jar_ro, GFP_KERNEL);
        if (!new_ro)
            panic("override_creds(): kmem_cache_alloc() failed");

        use_cnt_ptr = kmem_cache_alloc(usecnt_jar,GFP_KERNEL);
        if(!use_cnt_ptr)
            panic("override_creds() : Unable to allocate usage pointer\n");

        tsec = kmem_cache_alloc(tsec_jar, GFP_KERNEL);
        if(!tsec)
            panic("override_creds() : Unable to allocate security pointer\n");
    }

    rkp_cred_fill_params(new,new_ro,use_cnt_ptr,tsec,RKP_CMD_OVRD_CREDS,rkp_use_count);
    rkp_call(RKP_CMDID(0x46),(unsigned long long)&cred_param,0,0,0,0);

    rocred_uc_set(new_ro,2);
    rcu_assign_pointer(current->cred, new_ro);
#endif /* CONFIG_RKP_KDP */
}
```

- To override process's credentials

- Allocate new cred from RO kmem\_cache

- Ask RKP to update current cred and security context

# Further cred verifying in secure world

- rpk\_override\_creds()
  - -> rkp\_call(RPK\_CMD(0x41)) -> rkp\_assign\_creds()
- rkp\_assign\_creds()
  - Real implementation of override\_cred() in secure world
  - Additional verifying in KNOX 2.6

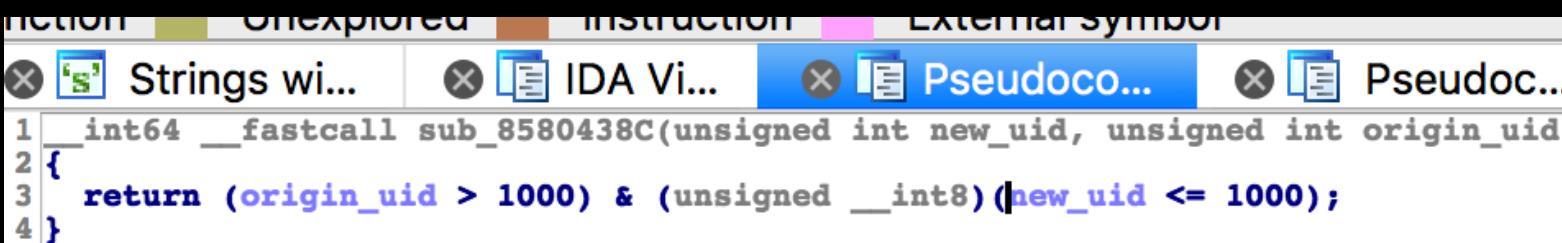
```
if ( !v1 )
    return rkp_printh("NULL pData", OLL, OLL, OLL);
_memcpy(&priv_escalcation, v1, sizeof(v1));
result = integrity_chk();
if ( result )
{
    v3 = get_caller_thread_info();
    pcb = get_physical_addr(v3);
    old_cred = get_physical_addr(*(_QWORD *) (pcb + 8 * hardcoded_table[17]));
    new_cred = get_physical_addr(v1);
    v7 = get_physical_addr(v12);
    new_cred_copy = v7;
    if ( new_cred && v7 && !sub_85803838(v7) )
    {
        if ( (unsigned int)check_there_is_adbd_zygote(pcb, old_cred) && (unsigned int)uid_checking(new_cred, old_cred) )
        {
            rkp_printh(
                "Priv Escalation!",
                new_cred,
                *(_QWORD *) (new_cred + 8LL * HIDWORD(hardcoded_table[19])),
                *(_QWORD *) (old_cred + 8LL * HIDWORD(hardcoded_table[19])));
            result = priv_escalation_abort(new_cred, old_cred, 1LL);
        }
    }
}
```

• Part of Data flow integrity

• UID checking

# uid\_checking()

- Check if adbd and zygote has started up
  - If not, allow the override
  - If true, the Android initialize has been finished, start UID checking
- Unprivileged process(uid>1000) cannot override the credential with high privilege(uid 0~1000)
- But still can change its kernel capabilities (very important!)



The screenshot shows the IDA Pro interface with the "Pseudocode" tab selected. The assembly code is as follows:

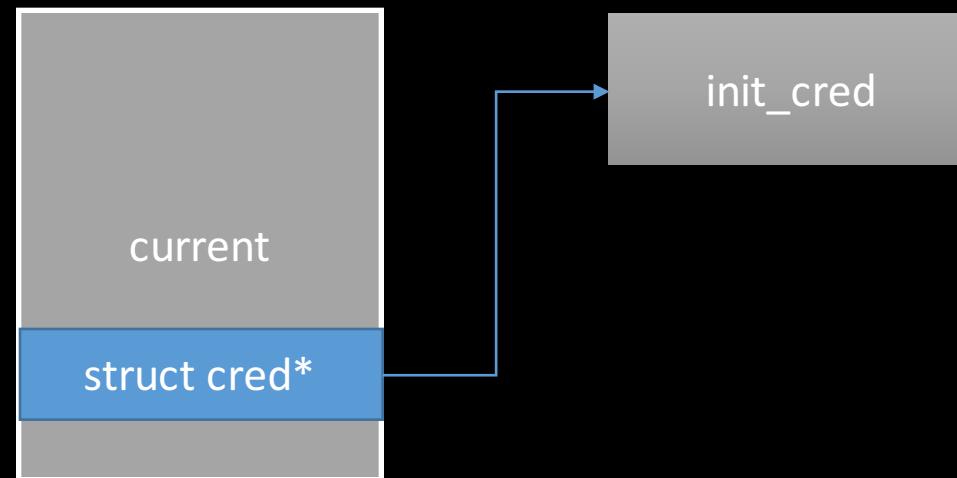
```
1 __int64 __fastcall sub_8580438C(unsigned int new_uid, unsigned int origin_uid)
2 {
3     return (origin_uid > 1000) & (unsigned __int8)(new_uid <= 1000);
4 }
```

# integrity\_checking()

- Do similar checking with security\_integrity\_current() in Linux Kernel
- Will analyze security\_integrity\_current() later

# Another old trick to change credential

- For now we know credentials are READ-ONLY
- What if we reuse init's credential?



- Not working on S7,because of Data Flow Integrity

# Data Flow Integrity

- New in KNOX 2.6
- Implemented in both Linux kernel and Secure world
  - security\_integrity\_current() (kernel)
  - Integrity\_checking()
- Additional members in struct cred{}

```
#endif
    struct user_struct *user; /
    struct user_namespace *user_n
    struct group_info *group_info
    struct rcu_head rru; /
#endif CONFIG_RKP_KDP
    atomic_t *use_cnt;
    struct task_struct *bp_task;
    void *bp_pgd;
    unsigned long long type;
#endif /*CONFIG_RKP_KDP*/
};

#endif CONFIG_RKP_KDP
```

- use\_cnt: pointer to refcount of cred
- bp\_task: pointer to this cred's owner
- bp\_pgd: pointer to process's PGD
- type: not used in DFI

# Data Flow Integrity

- New in KNOX 2.6
- Implemented in both Linux kernel and Secure world
  - security\_integrity\_current() (kernel)
  - Integrity\_checking()
- Additional members in struct task\_security\_struct{}

```
struct task_security_struct {  
    u32 osid;          /* SID prior to last exec */  
    u32 sid;           /* current SID */  
    u32 exec_sid;      /* exec SID */  
    u32 create_sid;    /* fscreate SID */  
    u32 keycreate_sid; /* keycreate SID */  
    u32 sockcreate_sid; /* fscreate SID */  
#ifdef CONFIG_RKP_KDP  
    void *bp_cred;  
#endif  
};
```

- bp\_cred: pointer to this context's owner cred

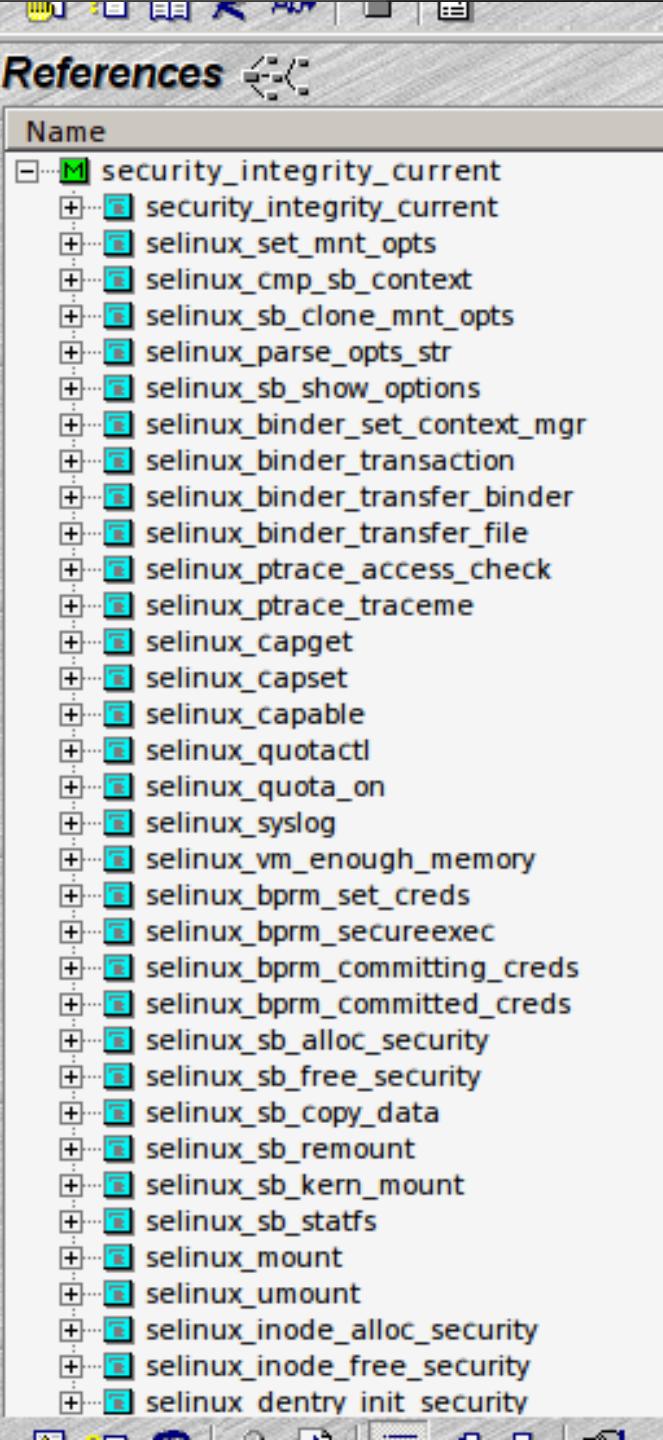
# security\_integrity\_current()

- Hard-coded hooking in every SELinux routines
  - Verify process's credential in real-time
  - To check if
    - current *struct cred{}* and *struct task\_security\_struct{}* are allocated in RO page
    - cred->bp\_task is current process's
    - task\_security->bp\_cred is current cred
    - current mount namespace is malformed

```
}

/* Main function to verify cred security context of a process */
int security_integrity_current(void)
{
    if ( rkp_cred_enable &&
        (rkp_is_valid_cred_sp((u64)current_cred(),(u64)current_cred()->secu
        cmp_sec_integrity(current_cred(),current->mm)) ||
        cmp_ns_integrity())))
    {
        rkp_print_debug();
        panic("RKP CRED PROTECTION VIOLATION\n");
    }
    return 0;
}

unsigned int rkp_get_task_sec_size(void)
```



# Summary of RKP and DFI

- Even we achieved arbitrary kernel memory overwriting, we cannot:
  - Manipulate credentials and security context in kernel mode
  - Point current credential to init\_cred
  - Call rkp\_override\_creds() to ask secure world to help us override credential with uid 0~1000
- But we still can:
  - Call kernel function from user mode
    - Hijacking *ptmx\_fops->check\_flags(int flag)*
    - The number of parameters is limited
    - Only low 32bit of X0 is controllable
  - Override credential with full kernel capabilities (*cred->cap\_\*\**)
  - Overwrite unprotected data in kernel

# Bypassing RKP and DFI

- Main idea: ask kernel to create a privileged process for me
- Creating a root process
- I can't call *call\_usermodehelper(path, argv, envp, wait)* via *ptmx\_fops->check\_flags(flag)*
- Call *orderly\_poweroff()* instead

# orderly\_poweroff()

```
/**  
 * orderly_poweroff - Trigger an orderly system poweroff  
 * @force: force poweroff if command execution fails  
 *  
 * This may be called from any context to trigger a system shutdown.  
 * If the orderly shutdown fails, it will force an immediate shutdown  
 */  
int orderly_poweroff(bool force)  
{  
    if (force) /* do not override the pending "true" */  
        poweroff_force = true;  
    schedule_work(&poweroff_work);  
    return 0;  
}  
EXPORT_SYMBOL_GPL(orderly_poweroff);
```

- Call `__orderly_poweroff()` in worker thread

```
char poweroff_cmd[POWEROFF_CMD_PATH_LEN] = "/sbin/poweroff";  
  
static int __orderly_poweroff(bool force)  
{  
    char **argv;  
    static char *envp[] = {  
        "HOME=/",  
        "PATH=/sbin:/bin:/usr/sbin:/usr/bin",  
        NULL  
    };  
    int ret;  
  
    argv = argv_split(GFP_KERNEL, poweroff_cmd, NULL);  
    if (argv) {  
        ret = call_usermodehelper(argv[0], argv, envp, UMH_WAIT_EXEC);  
        argv_free(argv);  
    } else {
```

- Worker create a new root process, cmd is `poweroff_cmd`
- `poweroff_cmd` is writeable

# Bypassing steps

- Call `rpk_override_creds()` via `ptmx_fops->check_flags()`
  - Override own cred to gain full kernel capabilities
  - But don't change uid
- Overwrite `poweroff_cmd` with “/data/data/\*\*/ss7kiler”
- Call `orderly_poweroff()` via `ptmx_fops->check_flags()`
- Modify ss7killer's `thread_info->address limit`
- ss7killer: call `rpk_override_creds()` to change its sid from `u:r:kernel:s0` to `u:r:init:s0`

# Result: privileged ss7killer

- root
- u:r:init:s0

```
[shell@hero2qltechn:/ $ ps -Z | grep init
u:r:init:s0          root      1      0      /init
u:r:init:s0          root    20592  1      /data/local/tmp/ss7killer
u:r:init:s0          root    20608  20592  su
u:r:init:s0          root    20611  1      daemonsu:mount:master
u:r:init:s0          root    20614  1      daemonsu:master
u:r:init:s0          root    20797  20614  daemonsu:0
u:r:init:s0          root   21161  20614  daemonsu:10193
u:r:init:s0          root   21163  21161  daemonsu:10193:21157
```

# u:r:init:s0

- Not good enough
- Still be limited by SELinux
- Almost can do nothing...
- Disabling/Bypassing SELinux is necessary

# SELinux enhancement

- Disabled CONFIG\_SECURITY\_SELINUX DEVELOP long time ago
  - Cannot disable SELinux by overwrite selinux\_enforcing
  - Statically enforcing all the time
- init process cannot reload SELinux policy after system initialized
- Permissive domain is not allowed

# Permissive domain

- Officially used by Google before Lollipop
  - For policy developing purpose
- All domains are non-permissive since Lollipop
- Domains still can be switched to permissive mode by policy reloading (/sys/fs/selinux/load)

**AndroidXRef** KitKat 4.4.4\_r1

xref: /external/sepolicy/init.te

---

[Home](#) | [History](#) | [Annotate](#) | [Line#](#) | [Navigate](#) | [Download](#)

```
1 # init switches to init domain (via init.rc).
2 type init, domain;
3 permissive init;
4 # init is unconfined.
5 unconfined_domain(init)
6 tmmfs_domain(init)
```

# Permissive domain – kernel support

- A permissive domain's access vector decision(AVD) will be set AVD\_FLAGS\_PERMISSIVE
- All operations are permitted

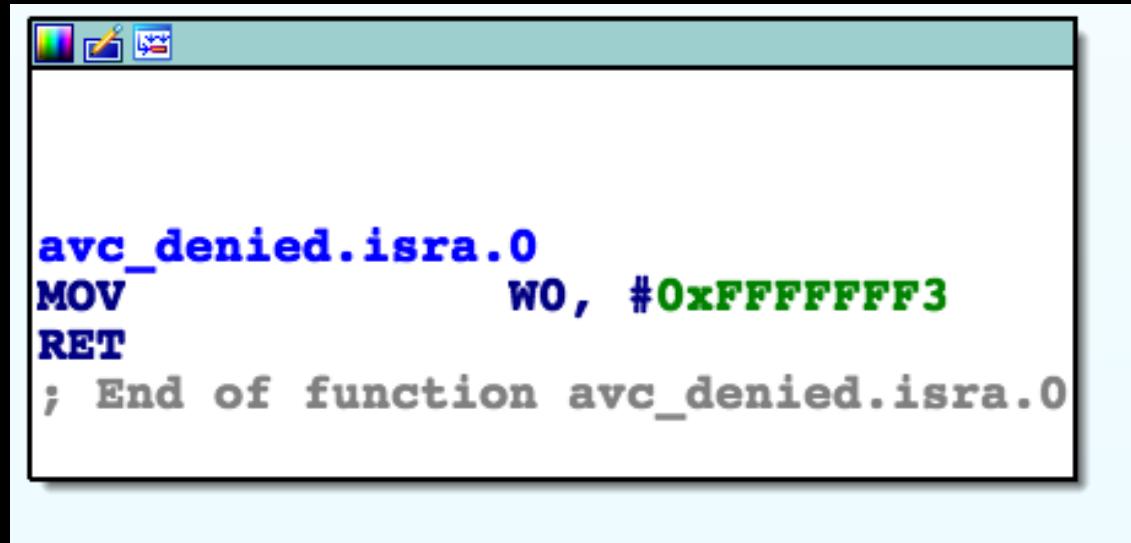
```
static __attribute__((noinline)) int avc_denied(u32 ssid, u32 tsid,
                                             u16 tclass, u32 requested,
                                             u8 driver, u8 xperm, unsigned flags,
                                             struct av_decision *avd)
{
    if (flags & AVC_STRICT)
        return -EACCES;

    if (selinux_enforcing && !(avd->flags & AVD_FLAGS_PERMISSIVE))
        return -EACCES;

    avc_update_node(AVC_CALLBACK_GRANT, requested, driver, xperm, ssid,
                    tsid, tclass, avd->seqno, NULL, flags);
    return 0;
}
```

# S7 removed AVD\_FLAGS\_PERMISSIVE

- `avc_denied` always simply return -EACCES



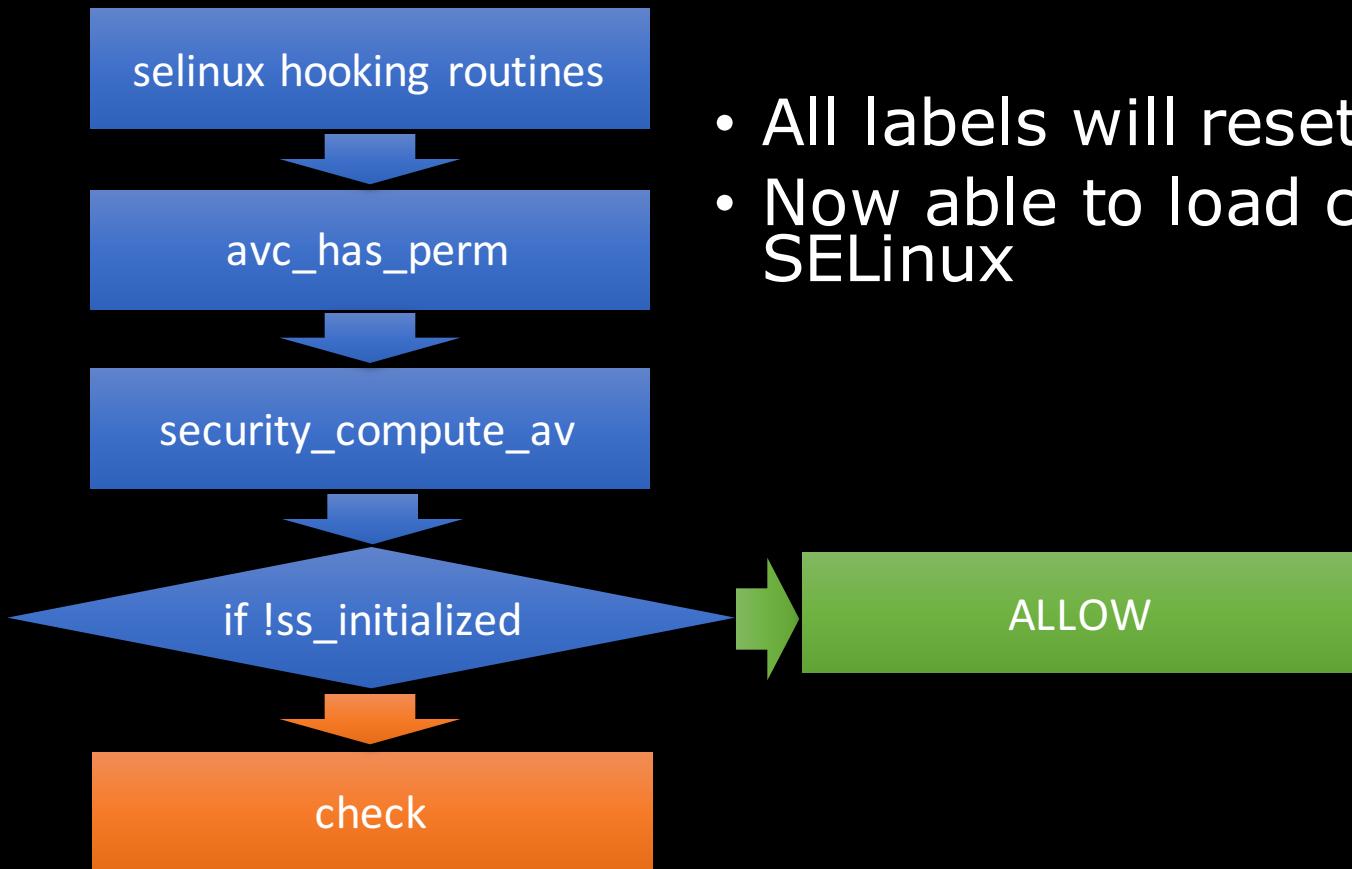
A screenshot of a debugger interface showing assembly code. The window has a title bar with icons for file, edit, and help. The assembly code is as follows:

```
avc_denied.isra.0
MOV             W0, #0xFFFFFFFF3
RET
; End of function avc_denied.isra.0
```

# Bypass SELinux on S7

Tencent

- Cheating kernel that SELinux is not initialized yet
  - Depends on global variable `ss_initialized` (writable)



- All labels will reset to none except kernel domain
- Now able to load customized policy and reinitialize SELinux

# After setting ss\_initialized = 0

- All labels missed except kernel
  - SELinux must be re-enabled ASAP, or Apps may corrupt files' label permanently
- Load customized policy and reinitialize SELinux

```
kernel          root    900   2      dhd_watchdog_th
kernel          root    901   2      dhd_rpm_state_t
-              system  903   1      /system/bin/factory.adsp
-              net_admin 910   1      /system/bin/ipacm
-              radio    914   1      /system/vendor/bin/qti
-              radio    922   1      /system/bin/netmgrd
-              radio    947   1      /system/bin/rild
-              jack     950  694      androidshmservice
kernel          root   1157   2      irq/140-arm-smm
kernel          root   1160   2      irq/141-arm-smm
-              system  1161   1      /system/bin/mcDriverDaemon
kernel          root   1209   2      tee_scheduler
kernel          root   1218   2      irq/162-arm-smm
kernel          root   1219   2      irq/167-arm-smm
kernel          root   1221   2      irq/163-arm-smm
kernel          root   1223   2      irq/168-arm-smm
-              system  1250  728      system_server
```

# Policy customizations

- Policy database locate at /sys/fs/selinux/policy
- Modify the database with libsepol API
  - Load policy DB to the user memory
  - Add rules into database
    - Allow untrusted\_app, init, toolbox domain to do everything
  - Ask kernel to reload the database
- Set ss\_initialized to 1

# Gain Root

- Leaking kernel information ✓
- Bypassing KASLR ✓
- Overwriting arbitrary memory ✓
- Bypassing RKP & DFI ✓
- Bypassing enforced SELinux ✓

The screenshot shows a terminal window with the following session:

```
Last login: Wed Jun 29 19:23:54 on ttys000
[retme-no-RMBP:~ ddshen$ adb shell
[shell@hero2qltechn:/ $ ps -Z | grep s7kill
u:r:init:s0                      root      14453 1907 /data/local/tmp/s7killer
[shell@hero2qltechn:/ $ ps -Z | grep daemon-su
u:r:init:s0                      root      14593 1     daemon-su:mount:master
u:r:init:s0                      root      14596 1     daemon-su:master
u:r:init:s0                      root      14620 14596 daemon-su:0
[shell@hero2qltechn:/ $ su
[root@hero2qltechn:/ # id
uid=0(root) gid=0(root) groups=0(root) context=u:r:toolbox:s0
[root@hero2qltechn:/ # getenforce
Enforcing
root@hero2qltechn:/ #
```

A notification bubble in the top right corner of the window says "有更新项目 您要立刻重新启动以安试今晚安装?" (There are updates available. Do you want to restart now to install them?).

Tencent

